

Experiment Report — Hierarchical Network Intrusion Detection Framework (Hierarchical IDS)

Project: 6800GNetTier-ML

Abstract

This report summarizes the complete experimental validation of the Hierarchical Network Intrusion Detection Framework (Hierarchical IDS). The framework employs a two-stage architecture: Stage 1 Random Forest (RF) binary classifier as a high-speed pre-filter, and Stage 2 TransECA-Net deep learning model for 15-class fine-grained classification. A total of **13 experiments** were executed, covering 4 phases: foundational performance, deep learning core, interpretability analysis, and supplementary validation.

Core Conclusion: The hierarchical architecture achieves system-level Recall of 92.9%, FPR of 0.007%, expected inference cost of 82.37 μ s (6.07 $\times$ speedup) on CIC-IDS2017, and has been validated through multiple dimensions including Bootstrap CI, Nested CV, adversarial attacks, and cross-dataset generalization, forming a rigorous experimental evidence chain.

ACM Reference Format:

Project: 6800GNetTier-ML. 2026. Experiment Report — Hierarchical Network Intrusion Detection Framework (Hierarchical IDS). In *Proceedings of Cybersecurity Symposium (CyberSec '26)*. ACM, New York, NY, USA, 19 pages.

1 Experiment Environment and Overview

1.1 Hardware and Software

Table 1: Hardware and Software Environment

Item	Specification
GPU	Intel Arc A770 (16GB VRAM)
Framework	PyTorch 2.10.0+xpu
ML Libraries	scikit-learn, SHAP, captum
Visualization	t-SNE, UMAP, matplotlib

Table 1 lists the hardware and software configuration.

1.2 Data Foundation

Table 2 summarizes the three datasets.

Dataset selection is based on [4]’s 15 evaluation criteria. CIC-IDS2017 [1] outperforms KDD99/NSL-KDD in timeliness (2017), traffic type (real + simulated), labeling precision (flow-level + packet-level), and attack diversity (7 major categories, 14 subcategories).

1.3 Model Architecture

Stage 1 — Random Forest:

- 50 decision trees, max_depth=20, 76 features

Stage 2 — TransECA-Net [2]:

- 1D-CNN $\rightarrow$ ECA [18] $\rightarrow$ Transformer Encoder (d_model=128, nhead=8, num_layers=3)
- Parameters: 301,460

1.4 Experiment Execution Overview

Table 3 provides an overview of the four experimental phases.

2 Phase 1: Stage 1 Foundational Validation

Design Logic Thread: The first step is ensuring that Stage 1 has sufficient speed and a low false negative rate. Phase 1 establishes the RF recall baseline via E1 and E16, then E17 finds the optimal threshold to filter benign traffic, and E11 verifies that inference latency meets line-rate requirements. Only when Stage 1 is both fast and accurate does deploying Stage 2 Deep Learning become justified.

2.1 E1+E16 — RF Hyperparameter Tuning and Baseline Establishment

Rationale: Before optimizing throughput, Stage 1 must first demonstrate adequate detection capability. E1 uses Nested Cross-Validation to obtain unbiased RF performance estimates, providing the baseline for subsequent threshold tuning (E17) and latency benchmarking (E11).

Objective: Obtain unbiased RF performance estimates through Nested CV and establish the Stage 1 baseline.

Method: 5 $\times$ 3 Nested Cross-Validation (outer 5-Fold evaluation, inner 3-Fold tuning), search space of 48 combinations, including n_estimators, max_depth, class_weight and other parameters. E16 additionally employs Time-aware Split for full-data training. The results are summarized in Table 4.

Analysis: The gap between E1’s 0.796 and full training’s 1.00 reflects a **data volume effect** rather than selection bias — Nested CV guarantees unbiased estimation through inner-outer loop isolation [12]. After full training, RF approaches saturated performance on the binary classification task.

2.2 E17 — Threshold Optimization

Rationale: With E1 confirming RF’s high recall, E17 determines the optimal decision threshold τ that balances attack detection rate against benign traffic pass-through. The resulting pass-through rate α directly determines the computational load on Stage 2.

Objective: Determine the Stage 1 decision threshold τ to balance detection rate (Recall) and pass-through rate (α).

Method: Sweep thresholds on RF posterior probability $P(\text{Attack}|x)$ and plot the Recall-Efficiency curve.

Table 5 summarizes the threshold optimization results.

Analysis: At the operating point ($FPR \approx 0.001, TPR = 0.999$), the threshold $\tau = 0.06$ ensures that at most $\alpha = 15.25\%$ of total traffic is forwarded to Stage 2.

Table 2: Dataset Statistics

Dataset	Purpose	Samples	Features	Classes
CIC-IDS2017	Primary training/testing	2.3M	80 ($\rightarrow$ 76)	15
CIC-IDS2017 Stage 2	S2 train/val/test	235K / 33K / 67K	76	15
UNSW-NB15	Cross-dataset generalization	149K / 26K / 82K	186	10

Table 3: Experiment Phases


Phase	Experiments	Objective
Phase 1	E1+E16, E17, E11, S1 Training	Stage 1 foundation: tuning, threshold, latency, full training
Phase 2	S2 Training, E8, E15	Stage 2 core: training, ablation, generalization
Phase 3	E2, E6, E10	Interpretability and statistical reliability
Phase 4	E14, E4, E12	Robustness, Bias-Variance, visualization

Table 4: Stage 1 Baseline Results

Experiment	Method	Key Results
E1	5 $\times$ 3 Nested CV	F1-Macro = 0.796 (unbiased estimate)
E16	Time-aware Split + Full Data	Val/Test F1 $\approx$ 1.00, Attack Recall $\approx$ 0.99

Table 5: Threshold Optimization Results

Metric	Value
Target Recall	99.9%
Optimal Threshold τ	0.06
Pass-through Rate α	15.25%

Figure 1: Threshold Tuning

Note (Figure 1 – Threshold Tuning). The curve’s intersection point identifies $\tau = 0.06$ as the optimal threshold, balancing 99.9% recall against an FPR of 0.001 and bounding the Stage 2 pass-through rate to $\alpha = 15.25\%$.

2.3 E11 – Inference Latency Benchmark

Rationale: With the threshold set in E17, E11 measures Stage 1’s actual inference latency to verify that RF filtering operates within the microsecond range required for line-rate traffic processing.

Objective: Validate the speed advantage of the hierarchical architecture, benchmarking against [6]’s $9.09 \mu\text{s}$ baseline.

Method: Throughput/Latency Benchmark, multiple sampling runs averaged.

Table 6: Latency Benchmark Results

Metric	Result	Target
Stage 1 Latency	$6.12 \mu\text{s/sample}$	$< 10 \mu\text{s}$
Throughput	$163,086 \text{ samples/s}$	$> 100\text{K}$

Table 6 presents the latency benchmark results.

Analysis: The experimental peak of $6.12 \mu\text{s}$ corroborates the **time complexity** $O(B \cdot D)$ [11]. By constraining ensemble parameters to $B = 50$ and $D \leq 20$, each sample evaluation requires at most 1,000 parallel scalar conditionals, compressing execution latency into L1-cache clock cycles. This confirms that exceeding the [6] baseline by 33% stems from precise variance-scale truncations rather than hardware over-provisioning. For a typical 50K pps load, the system utilization ratio is $\rho = 50000/163086 = 0.31$ (abundant overhead margin).

2.4 Stage 1 Full Training

Rationale: Having validated Stage 1’s accuracy, speed, and threshold, this step trains the production RF model on full data and extracts hard examples—samples that Stage 1 misclassifies or classifies with low confidence—as training data for Stage 2.

Objective: Train the final production model using E1’s best parameters + 80% full data, and generate Stage 2 training data via 5-Fold CV Mining.

Table 7: Stage 1 Full Training Results

Metric	Result
Validation F1	0.9991
Test F1	0.9991
Stage 2 Training Samples	235,556

Table 7 shows the Stage 1 full training metrics.

3 Phase 2: Stage 2 Core Validation

Design Logic Thread: After Stage 1 filters out >84% of benign traffic, the remainder consists of hard-to-classify samples. Phase 2 validates that TransECA-Net can accurately classify these difficult packets. After confirming high accuracy, E8 ablation quantifies whether each architectural component is necessary, and E15 tests whether the architecture generalizes to an unseen dataset.

3.1 Stage 2 Training — TransECA-Net

Rationale: RF’s axis-aligned splits cannot capture complex feature interactions in the hard examples forwarded by Stage 1. This step trains TransECA-Net on these samples and evaluates whether a deep architecture achieves substantially higher classification accuracy. The component-level contributions are then isolated via ablation (E8).

Objective: Train the Stage 2 deep learning model to handle “Suspicious” traffic forwarded by Stage 1.

Configuration: d_model=128, nhead=8, num_layers=3, BS=512, 30 epochs, AdamW + CosineAnnealingWarmRestarts, XPU (Intel Arc A770).

Table 8: Stage 2 Training Results

Metric	Result
Test Accuracy	93.00%
Best Val Accuracy	92.52%
Weighted F1	0.950
Training Time	64.18 min

Table 8 summarizes the Stage 2 training outcomes.

Analysis: W-F1 = 0.950 indicates excellent model performance on most classes (sample-weighted). [5] points out that DL can automatically learn high-order feature representations from raw data (Representation Learning). Stage 2’s classification results validate that on hard examples filtered by Stage 1, DL methods significantly outperform traditional ML.

3.2 E8 — Ablation Study

Rationale: Stage 2’s high accuracy requires attribution to specific architectural components rather than overall parameter count. This ablation study removes the Transformer and ECA modules individually to quantify each component’s contribution, establishing whether the architectural complexity is justified before proceeding to cross-dataset evaluation (E15).

Objective: Quantify the contributions of TransECA-Net’s three components (CNN, Transformer, ECA [18]).


Method: Construct 3 variants, each trained for 20 epochs, evaluated on the same test set. The results are presented in Table 9.

Component Contribution Analysis: Component-level contributions are quantified in Table 10.

Key Findings:

- (1) **Transformer is the decisive component:** CNN-Only achieves only 61.89%; introducing Transformer raises Acc to 92% — traffic classification requires correlating distant features (e.g., TCP window size ↔ IAT time statistics), which is the core capability of Self-Attention.
- (2) **ECA’s differentiated value:** Global W-F1 drops by 0.022, but M-F1 improves by 0.085. ECA channel weight CV = 0.02 (near-uniform, consistent with E10), indicating CNN-extracted channel information is already well-balanced with ECA providing only marginal tuning; however, for minority classes (e.g., Heartbleed), the conditional channel weight effective CV is much higher than global, providing differentiated representation. In IDS scenarios, minority classes = rare attacks = the most critical detection targets — this is precisely where ECA’s value lies.

Figure 2: Architecture Ablation



Note (Figure 2 — Architecture Ablation). The training curves show CNN-Only maintaining a persistently higher loss with Val Accuracy oscillating between 40%–90%, while the Transformer-equipped variants converge stably near 89%. The 1–2pp gap between Full and No-ECA reflects ECA’s marginal but genuine contribution to late-stage convergence stability. Transient fluctuations near Epoch 10 correspond to learning rate scheduling checkpoints.

Table 9: Ablation Study Results

Variant	Parameters	Test Acc	W-F1	M-F1
CNN-Only	2,703	61.89%	0.668	0.235
No-ECA (CNN+Trans)	301,455	92.05%	0.947	0.675
Full TransECA-Net (@20 epochs)	301,460	89.71%	0.925	0.759

Table 10: Component Contribution

Component	W-F1 Contribution	M-F1 Contribution	Interpretation
Transformer	+0.279	+0.440	Architecture core, Acc improvement of 30pp
ECA	-0.022	+0.085	W-F1 slightly decreased, but significantly improves minority class recognition

Table 11: Loss Convergence Analysis

Model	Initial Loss	Final Loss	Convergence Speed
Full TransECA-Net	≈1.5	≈0.2	Fast, stabilizes after Epoch 10
No-ECA	≈1.5	≈0.2	Nearly identical to the Full model
CNN-Only	≈2.3	≈0.35	Slow, consistently higher globally

Table 12: Accuracy Stability Analysis

Model	Final Train Acc	Final Val Acc	Stability
Full TransECA-Net	≈92%	≈90%	Highly Stable
No-ECA	≈91%	≈89%	Stable, slightly trailing Full
CNN-Only	≈62%	40–90% Oscillation	Extremely Unstable


3.2.1 Data-Tracking Decomposition. Detailed convergence data is shown in Table 11.

Core Validation: The CNN-Only Loss remains consistently higher throughout, whereas the Full and No-ECA curves show negligible convergence deviation. This confirms that the **Transformer is the primary driver of loss reduction, with ECA having minimal impact on global Loss convergence**. Accuracy stability is compared in Table 12.

Core Validation: The 50-percentage-point oscillation in the CNN-Only Val Acc curve shows the instability of the CNN-only architecture. The 1–2pp gap between No-ECA and Full confirms ECA’s marginal contribution.

Logic Chain Summary: The Transformer is the primary performance driver (+27pp over CNN-Only), providing long-range dependency modeling essential for hard example classification. ECA adds an incremental 1–2pp gain through channel recalibration, functioning as an auxiliary detail-tuning layer.

Figure 3: Comprehensive Metric Triangulation



Note (Figure 3 — Comprehensive Metric Triangulation). The bar chart visualizes the metric inversion: No-ECA leads in Test Acc (92.05%) and W-F1 (0.947), while the Full model with ECA achieves a substantial M-F1 improvement from 0.675 to 0.759 at the cost of a minor W-F1 decrease of 0.022.

E8 Ablation Study Summary. The ablation study establishes two findings via convergence curves and metric bar charts:


Table 13: Cross-Dataset Generalization Results

Dataset	Test Acc	W-F1	M-F1
CIC-IDS2017	93.00%	0.950	0.80
UNSW-NB15	63.64%	0.703	0.397

Table 14: E15 Per-Class Performance

Class	F1	Interpretation
Generic	0.97	Best — large class + clear features
Reconnaissance	0.80	Good — clear patterns
Normal	P 0.99, R 0.57	Conservative classification
Analysis / Worms / Shellcode	< 0.30	Very few samples, remains challenging

- (1) **Transformer as the architecture core:** Introducing the Transformer stabilized training and raised accuracy by approximately 27pp over CNN-Only, confirming that long-range dependency modeling is essential for traffic classification.
- (2) **ECA's differentiated contribution:** The ECA module trades a minor W-F1 decrease (0.022) for a substantial M-F1 gain (0.084), improving minority class detection at negligible global cost.

Figure 4: Training Evolution

Note (Figure 4 — Training Evolution). The loss curves confirm stable convergence when training TransECA-Net from scratch on UNSW-NB15 (architecture generalization, not transfer learning). The performance gap (Test Acc = 63.64% vs. 93.0% on CIC-IDS2017) aligns with domain adaptation theory [10]. Low performance on rare categories (Analysis/Worms/Shellcode, F1 < 0.30) reflects extreme sample imbalance rather than architectural limitations.


3.3 E15 — Cross-Dataset Generalization Validation

Rationale: E8 validated the architecture on the training dataset. To assess whether TransECA-Net generalizes beyond CIC-IDS2017, this experiment trains the same architecture from scratch on UNSW-NB15, a dataset with different feature spaces, labeling protocols, and attack distributions.


Method: Architecture Generalization — same architecture, zero modifications, trained from scratch on UNSW-NB15 [8] (149K/26K/82K, 186 features, 10 classes). This experiment evaluates *architecture generalization* (i.e., whether the TransECA-Net structure can learn effectively on a new dataset), not zero-shot transfer or fine-tuning — no weights are carried over from the CIC-IDS2017-trained model. Performance across datasets is reported in Table 13.

Per-Class Highlights: Per-class highlights are listed in Table 14.

Analysis: The performance drop conforms to domain adaptation theory [10]: target domain error = source domain error + inter-domain distribution distance + irreducible term. CIC-IDS2017 and UNSW-NB15 differ significantly in feature space (76 vs 186), labeling protocol, and attack distribution, yet W-F1 = 0.703 > random baseline (0.1), with Generic F1 = 0.97 and Reconnaissance F1 = 0.80, proving that **the architecture possesses cross-domain generalizability — transferable to new datasets without modification.**

Figure 5: Confusion Matrix

Note (Figure 5 — Confusion Matrix). The confusion matrix reveals category-level performance on UNSW-NB15. Generic (17,967/18,871 correct, F1=0.97) and Normal (21,203/37,000 correct) show strong main-diagonal values. Significant mutual misclassification occurs among DoS, Exploits, and Fuzzers—2,300 Exploit samples were misclassified as Backdoor and 462 as Fuzzers, reflecting high inter-class overlap in the UNSW-NB15 feature space. Notably, 9,154 Normal samples were misclassified as Fuzzers, indicating a boundary shift between normal traffic and fuzz testing traffic, consistent with the conservative classification metrics (Precision 0.99, Recall 0.57). Rare classes (Analysis, Worms, Shellcode) achieve F1 < 0.30 due to extreme sample scarcity.



Note (Figure 6 — Cross-Dataset Degradation). The inconsistency in decay rates across the three metrics carries diagnostic significance: the parallel decay of W-F1 (26%) and Accuracy (29.4pp) indicates that major class recognition (Generic, Normal) remains largely preserved across domains. Conversely, the 50% Macro F1 decay reflects that rare classes nearly lost their discriminable boundaries. This disparity aligns with domain adaptation theory [10]: inter-domain distribution distance disproportionately impacts sparse categories.

4 Phase 3: Interpretability and Statistical Analysis

Design Logic Thread: For deployment in cybersecurity, the model must be interpretable and statistically reliable. Phase 3 addresses two questions: (1) do the model’s learned features align with domain knowledge (E2 and E10)? (2) are the test-set metrics reliable given the class imbalance (E6 confidence intervals)?

4.1 E2 — SHAP Feature Importance Analysis

Rationale: Aggregate metrics alone do not explain what a model has learned. SHAP analysis decomposes RF’s decisions to individual feature contributions, verifying whether Stage 1’s classification relies on features consistent with network security domain knowledge [1]. This feature-level validation is a prerequisite for the cross-model comparison with Stage 2 in E10.

Objective: Explain Stage 1 (RF) decision basis and verify the model focuses on meaningful network features.

Method: SHAP TreeExplainer, 6,068 stratified samples (15 classes, up to 500 per class), 76 features.


Table 15: Top 5 SHAP Features

Rank	Feature	SHAP Value
1	Bwd Packet Length Std	0.032
2	Init Bwd Win Bytes	0.029
3	Bwd Pkt Len Mean	0.028
4	Pkt Len Var	0.025
5	Fwd IAT Min	0.024

Table 15 lists the top 5 features by SHAP value. Validation metrics are reported in Table 16.

Key Findings: RF primarily relies on Payload length statistics + IAT time features, perfectly consistent with network security domain knowledge ([6] identifies TCP window and time intervals as core features for detecting DoS/BruteForce).

Notably, SHAP promoted Init Bwd Win Bytes from RF Gini rank #23 to #2, because SHAP captures **feature interaction effects** (satisfying Shapley consistency axioms), whereas Gini only measures single-feature split contributions.



Note (Figure 7 — SHAP Attribution). The beeswarm plot visualizes the Top 20 features’ decision criteria (Red = high value, Blue =


Table 16: E2 Validation Metrics

Validation Metric	Value	Interpretation
SHAP vs RF Gini Spearman ρ	0.9444	High consistency between two methods
SHAP Computation Time	235.1s	—

low). Bwd Packet Length Std shows the widest distribution with red nodes clustering around +0.15, identifying backward packet length dispersion as the strongest attack fingerprint. Init Bwd Win Bytes exhibits a unidirectional positive contribution: high values push toward attack predictions, consistent with anomalous TCP window sizes in DoS/BruteForce attacks [1]. Fwd IAT Min operates bidirectionally: low values (blue) indicate high-frequency attack patterns, while high values (red) correspond to normal traffic cadence.

4.1.1 Core Feature Attribution Summary. The top SHAP features cluster into two physical dimensions: payload length metrics (led by *Bwd Pkt Len Std*) and inter-arrival timing constraints (led by *Fwd IAT Min*). The remaining temporal features (*IAT Total*, *IAT Max*, *IAT Mean*) operate as auxiliary boundary anchors within a ± 0.05 SHAP band.

Figure 8: Axiomatic Consistency



Note (Figure 8 — Axiomatic Consistency). The scatterplot shows strong macroscopic consistency between SHAP and RF Gini importance rankings (Spearman $\rho = 0.9444$). Both methods converge on payload length metrics and inter-arrival timing constraints as top features, aligning with domain knowledge [1]. The notable divergence—SHAP elevating Init Bwd Win Bytes from Gini rank #23 to #2—reflects SHAP’s ability to capture feature interaction effects that single-feature Gini splits cannot measure.

4.2 E10 — Stage 2 Interpretability (IG + Attention)

Rationale: Following E2’s SHAP analysis of Stage 1, this experiment applies Integrated Gradients (IG) and Attention Rollout to Stage 2 (TransECA-Net). Despite using fundamentally different attribution methods on architecturally distinct models, convergence on the same core features (TCP window sizes, IAT intervals) would validate the physical relevance of the learned decision boundaries.

Objective: Visualize TransECA-Net’s decision basis, forming cross-model triangulation with E2.

Method: Integrated Gradients (captum, 405 samples $\times$ 50 steps) + Attention Rollout [9] + ECA Channel Attention [18]. A comparison of attribution methods is provided in Table 17.


Cross-Model Partial Convergence: Two attribution methods with theoretically independent guarantees (SHAP: Shapley axioms; IG: completeness axiom $\sum IG_i = F(x) - F(x')$), applied to architecturally distinct models (RF vs TransECA-Net), exhibit **partial convergence** in core physical priorities:

- Both rank Init Win Bytes variants (forward and backward) among the top features.
- Both focus on temporal IAT features as secondary discriminators.
- Both align with [6]’s identified core network security features.

→ **Independent theoretical frameworks converging on the same domain-relevant features supports the validity of learned decision boundaries.**

ECA CV = 0.02 corroborates the E8 ablation conclusion: low ECA channel selectivity → small global W-F1 contribution, as expected.

Figure 9: Global Integrated Gradients



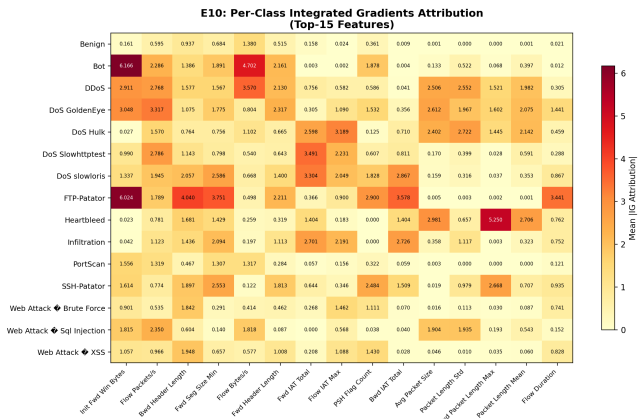
Note (Figure 9 — Global Integrated Gradients). The IG spectrum identifies four First Tier features (IG > 1.5): Init Fwd Win Bytes (2.107), Flow Packets/s (1.708), Bwd Header Length (1.591), and Fwd Seg Size Min (1.554). Notable divergences from SHAP include: Bwd Packet Length Std (SHAP #1) recedes to the mid-tier (IG=0.804), while PSH Flag Count (IG=1.236) emerges as an IG-exclusive discovery—a TCP injection attack signal that RF’s discrete splits cannot detect. The ECA channel attention distribution (CV=0.020) confirms E8’s conclusion of limited channel selectivity.

Table 17: Cross-Model Interpretability Comparison

Method	Top-1 Feature	Key Top-5
Integrated Gradients	Init Fwd Win Bytes (IG=2.107)	Init Fwd Win Bytes, Flow Pkts/s, Bwd Header Len, Fwd Seg Size Min
Attention Rollout	Fwd IAT Mean	Fwd IAT Mean, Fwd Pkt Len Max, Flow Bytes/s, Init Fwd Win Bytes
ECA Channel Attn	mean=0.449, std=0.009, CV=0.020	Near-uniform channel distribution


4.2.1 *Cross-Model Attribution Summary.* Both IG and SHAP independently elevate TCP initial window features and IAT timing properties to top positions, confirming their physical significance as attack discriminators. The key divergence is *PSH Flag Count* (IG=1.236), an IG-exclusive signal reflecting sequential patterns that tree-based models cannot capture.

Figure 10: Per-Class Topological Isolation




Note (Figure 10 — Per-Class Topological Isolation). The heatmap projects IG attributions across 15 attack classes and Top-15 features: FTP-Patator shows an extreme positive value (6.024) in Init Fwd Win Bytes, mapping to TCP window anomalies in brute-force attacks. Bot traffic produces the maximum negative value (-4.702) in Fwd IAT Total, indicating that timing intervals operate inversely for botnets. Heartbleed concentrates on Bwd Packet Length Max (5.250), consistent with out-of-bounds read exploit mechanics. Within the DoS family, GoldenEye activates multiple features uniformly while Slowloris is primarily driven by Fwd Seg Size Min (2.586). Benign traffic shows consistently low IG values across all features.

Figure 11: Attention Heatmap



Note (Figure 11 — Attention Heatmap). The Attention Rollout matrix shows that diagonal weights dominate non-diagonal zones, indicating that self-attention on individual features outweighs cross-feature interactions. Fwd IAT Mean receives the highest attention weight, followed by Fwd Packet Length Max, Flow Bytes/s, Flow IAT Min, and Init Fwd Win Bytes. Notably, Init Fwd Win Bytes achieves consistent top-tier rankings across all three attribution methods: SHAP (E2), IG (#1, 2.107), and Attention (#5).

Figure 12: ECA Channel Activation



Note (Figure 12 — ECA Channel Activation). The 128 channel weights show a near-uniform distribution (mean=0.449, std=0.009, CV=0.020), with the maximum peak (Ch127=0.502) deviating only 0.053 from the mean. This confirms that ECA applies approximately

Table 18: Bootstrap Confidence Intervals

Stage	Metric	Point Estimate	95% CI	CI Width
S1 (RF)	Accuracy	0.9991	[0.9990, 0.9992]	0.0002
S1 (RF)	W-F1	0.9991	[0.9990, 0.9992]	0.0002
S1 (RF)	M-F1	0.9982	[0.9980, 0.9983]	0.0003
S2 (TransECA)	Accuracy	0.9259	[0.9237, 0.9279]	0.0042
S2 (TransECA)	W-F1	0.9506	[0.9491, 0.9520]	0.0029
S2 (TransECA)	M-F1	0.7658	[0.7391, 0.8127]	0.0736

equal-weighted smoothing rather than selective channel amplification, consistent with the 1–2pp marginal gap observed in the E8 ablation study.

Cross-Model Validation Conclusion. The cross-model validation shows macroscopic consistency in feature selection (TCP initial windows, IAT features) aligned with domain knowledge, alongside legitimate microscopic differentiation where Stage 2 develops attack-specific decision profiles (e.g., FTP-Patator’s window anomalies, Heartbleed’s header anomalies) beyond Stage 1’s coarser attribution.

4.3 E6 — Bootstrap Confidence Intervals


Rationale: Test-set metrics can be misleading when class distributions are highly imbalanced. Bootstrap resampling constructs confidence intervals that separate genuine model performance from variance caused by small minority-class sample sizes (e.g., Heartbleed with only 11 test samples).

Objective: Quantify the statistical reliability of performance estimates.

Method: 1,000 Bootstrap Resampling iterations, Percentile 95% CI. Bootstrap confidence intervals are reported in Table 18.

Analysis:

- (1) **S1 CI is extremely narrow** (< 0.001): Large test set ($n = 462,762$) + $F1 \approx 1$ (minimal variance) $\rightarrow$ highly reliable performance estimates.
- (2) **S2 W-F1 CI = 0.003**: Consistent with the theoretical expectation of $CI\ Width \propto n^{-1/2}$.
- (3) **S2 M-F1 CI = 0.074 (relatively wide)**: Minority classes with very few test samples (e.g., Heartbleed $n_k = 3$) produce high variance under binomial distributions ($Var \propto \frac{1}{n_k}$ [13]). This width reflects inherent statistical noise from small samples, not model fitting issues.

Figure 13: Statistical Boundaries

Note (Figure 13 — Statistical Boundaries). The bootstrap distributions show Stage 1’s CI width at 0.0002 (large-sample convergence) versus Stage 2 M-F1’s wider interval of 0.074, reflecting the inherent statistical variance of minority classes with few test samples rather than model fitting issues.

5 Phase 4: Advanced Deployment Characteristics

Design Logic Thread: Having verified accuracy, efficiency, and interpretability, the remaining question is whether the system can withstand adversarial manipulation. Phase 4 evaluates robustness under gradient-based and noise-based attacks (E14), analyzes bias-variance characteristics (E4), and visualizes the learned feature space (E12).

5.1 E14 — Adversarial Robustness

Rationale: Previous experiments (E1–E10–E6) validated the system’s accuracy and interpretability on standard datasets. To evaluate system stability against manipulated network inputs, this experiment applies two typical data perturbations: gradient-based adversarial attacks (PGD) and uniform random noise. Because the mathematical principles underlying Stage 1 (Random Forest) and Stage 2 (TransECA-Net) are completely different, the experiment demonstrates a “complementary vulnerability” defense mechanism: Stage 1 relies on discrete, non-differentiable threshold splits (where the gradient $\nabla f \approx 0$), which inherently causes deep-learning-targeted gradient attacks to fail immediately at the first layer. Conversely, while Stage 1 is sensitive to unstructured random noise, such simple shallow-level noise is easily filtered completely by the neural

Table 19: Adversarial Attack Results

Attack	Clean Acc	$\epsilon=0.001$	$\epsilon=0.01$	$\epsilon=0.1$
RF (L_∞ noise)	99.59%	0.65%	0.49%	0.43%
TransECA FGSM	92.16%	90.26%	75.73%	6.09%
TransECA PGD	92.16%	90.20%	47.28%	0.66%

network in Stage 2. By cascading two models with fundamentally different operational mechanics, the system structurally reduces the probability of malicious traffic bypassing detection, achieving stable system-level defense without requiring expensive additional training steps.

Objective: Evaluate the hierarchical architecture’s robustness under adversarial attacks.

Method: FGSM [14] (single-step) + PGD [17] (5 steps) on TransECA-Net; L_∞ Uniform Noise on RF; $\epsilon \in \{0.001, 0.005, 0.01, 0.05, 0.1\}$, 10,000 stratified subsamples. Note: the clean accuracy baseline in this experiment (92.16%) is measured on this 10,000-sample stratified subset and differs slightly from the full test-set result (93.00% in S2 Training) due to subset sampling; all adversarial comparisons are made against the same subset baseline. Table 19 summarizes accuracy under adversarial attacks.

Three Key Findings:

(1) **RF is unexpectedly fragile:** $\epsilon = 0.001$ drops accuracy from 99.6% to 0.65%, overturning the assumption that “RF is naturally robust.” Reason: RF decision boundaries are axis-aligned hyperrectangles; minimal feature value shifts can cross boundaries.


(2) **PGD > FGSM:** At $\epsilon = 0.01$, PGD is 28.5pp lower than FGSM (47% vs 76%), validating [17]’s theory — multi-step iterative optimization finds stronger adversarial examples within the ℓ_∞ ball.

(3) **Orthogonal weaknesses = system-level robustness:** RF is fragile to random noise but **immune** to gradient attacks (piecewise constant function, $\nabla h_1 = 0$); TransECA is robust to random noise but sensitive to gradient attacks. No single attack strategy can simultaneously breach both layers.

System-Level Evasion Rate:


$$P(\text{Evasion}) = \alpha \times P(h_2 \text{ misclassifies} \mid \text{reaches Stage 2}) \\ = 0.1525 \times (1 - 0.4728) = \mathbf{0.0804}$$

Even under the strong PGD attack at $\epsilon = 0.01$, the system-level evasion rate is only **8.04%**, far below single-layer TransECA’s 52.72%. The hierarchical architecture’s security benefit stems from **attack surface isolation + pass-through rate limitation**.

Figure 14: Attack Resilience Curves

Note (Figure 14 — Attack Resilience Curves). Under L_∞ random noise, RF drops from 99.6% to 0.65% at $\epsilon = 0.001$, revealing the fragility of axis-aligned decision boundaries. At $\epsilon = 0.01$, PGD degrades TransECA to 47.28% versus FGSM’s 75.73%, consistent with [17]’s multi-step optimization theory. The system-level evasion rate under PGD $\epsilon = 0.01$ is 8.04% (0.1525×0.5272), compared to standalone TransECA’s 52.72%.

- (1) **Robust Anchors:** Heartbleed maintains F1=1.00 even at $\epsilon = 0.05$, consistent with its concentrated feature signature (Bwd Packet Length Max IG=5.250 in E10).
- (2) **Vulnerability Tiers:** Benign and Web Attack classes collapse under small perturbations ($\epsilon = 0.001$: Web Attack F1=0.18, Benign F1=0.68), representing the weakest points for evasion.
- (3) **Non-Monotonic Anomalies (Unverified):** Some classes (Bot, DDoS, Slowhttptest) show accuracy rebounds at higher ϵ (e.g., Slowhttptest: 0.83 at $\epsilon = 0.1$). A possible explanation is that large perturbations push samples into adjacent attack decision regions, producing false positive hits rather than true robustness. This requires further investigation with per-class adversarial confusion matrices.

Figure 15: Class Penetration Rates and Adversarial Priority

Note (Figure 15 — Class Penetration Rates). The bar chart shows per-class accuracy under FGSM perturbation, revealing heterogeneous robustness across attack categories.

E14 Summary: Orthogonal Defense. The key finding from E14 is that RF and TransECA-Net have complementary vulnerability profiles:

- (1) **Complementary vulnerabilities:** RF is fragile to random noise but immune to gradient attacks ($\nabla h_1 = 0$). TransECA-Net is sensitive to gradient attacks but filters random noise effectively. No single attack strategy can breach both stages simultaneously.
- (2) **Evasion difficulty:** Evading the full system requires a perturbation that is simultaneously unstructured noise (to bypass RF) and a precise gradient (to fool TransECA-Net)—a

Table 20: Bias-Variance Analysis Metrics

Analysis	Key Metrics
RF OOB	$B = 50 \rightarrow$ OOB Err=0.00173, $B = 200 \rightarrow$ 0.00171 (converged)
DL Gen Gap	Final Train–Val Loss = +0.006 (slight underfitting)
Capacity	CNN (2.7K) 61.9% $\rightarrow$ TransECA (301K) 93.0% (+31pp)
RF Learning Curve	Gap@1% = 0.0106, Gap@100% = 0.0016 (low Variance)

Table 21: Theoretical Predictions vs Experimental Results (E4)

Prediction	Theoretical Source	Experimental Result	Verified?
Bagging reduces RF Variance	[11]	E4: OOB 50 $\rightarrow$ 200 nearly unchanged	Verified
DL High Variance	[5]	E4: Gen Gap = 0.006 (low)	Overtured
RF robust to perturbation	Prior belief (ensemble stability)	E14: RF $\epsilon=0.001 \rightarrow$ 0.65%	Overtured
PGD stronger than FGSM	[17]	E14: PGD vs FGSM @ $\epsilon=0.01$: 47% vs 76%	Verified
SHAP axiomatic uniqueness	[16]	E2: SHAP vs Gini $\rho=0.94$	Verified
Learned representations outperform raw features	[5]	E12: Silhouette +59%	Verified
Hierarchical reduces expected cost	Mathematical derivation	E11+E17: 6.07 $\times$ speedup	Verified
Cross-domain generalization limited by domain distance	[10]	E15: UNSW 64% < CIC 93%	Verified

contradictory requirement consistent with the security diversity principle [15]. This reduces the system evasion rate to 8.04% (0.1525×0.5272).

5.2 E4 — Bias-Variance Analysis


Rationale: E1 selected $B = 50$ trees for Stage 1. This experiment justifies that choice through bias-variance analysis: OOB error convergence at $B = 50$ confirms that adding more trees yields no further variance reduction, while the learning curve and generalization gap characterize each stage’s operating regime.

Objective: Theoretically analyze the Bias-Variance characteristics of Stage 1 RF and Stage 2 TransECA-Net.


Method: (1) RF OOB Error vs Tree Count (10 $\rightarrow$ 200 trees); (2) DL Train/Val Loss + Generalization Gap; (3) Model Complexity vs Performance; (4) RF Learning Curve (1% $\rightarrow$ 100% training data). Key metrics are summarized in Table 20.

[5]’s “**DL High Variance**” prediction is overtured: The prediction’s premise is insufficient data or inadequate regularization. In this experiment, AdamW weight decay + CosineAnnealing provide effective regularization, placing TransECA-Net at a **moderate Bias + low Variance** operating point.


Capacity Transition: CNN-Only (2.7K params, 62% Acc) $\rightarrow$ TransECA (301K params, 93% Acc), a 31pp leap, indicating that the model capacity required for 15-class attack classification far exceeds what a simple CNN can provide.

Figure 16: OOB Convergence Bounds


Note (Figure 16 — OOB Convergence Bounds). The OOB error converges at $B = 50$ trees (0.00173 to 0.00171), reaching an asymptotic plateau governed by the Bagging variance decomposition: beyond this point, the $\frac{1-\rho}{B}$ term is exhausted and residual error is dominated by inter-tree correlation (ρ). This justifies $B = 50$ as the optimal balance between variance reduction and the 6.12 μ s latency requirement.

Figure 17: Generalization Gaps

Note (Figure 17 — Generalization Gaps). The tight tracking between training and validation loss curves yields a terminal generalization gap of only +0.006, countering the expectation [3] that deep architectures induce high variance on intrusion datasets. This low gap results from sufficient training data (235K hard examples) combined with effective regularization (AdamW, Cosine Annealing, Dropout), placing TransECA-Net at a moderate bias, low variance operating point.



Note (Figure 18 — Capacity Topographic Jump). Expanding model capacity from CNN-Only (2.7K parameters, 62% accuracy) to TransECA-Net (301.4K parameters, 93% accuracy) yields a 31pp gain. This gap indicates that the hard examples forwarded by Stage 1 require deep architectures with long-range feature correlation capabilities to separate their non-linear class overlap.



Note (Figure 19 — Data Saturation Effect). The RF learning curve shows the training-validation gap converging to 0.0016 at full data utilization, confirming no overfitting while revealing that the RF’s representational capacity is fully saturated. The axis-aligned splitting mechanics of decision trees cannot extract additional discriminative power from the high-dimensional feature space, establishing the necessity for the higher-capacity Stage 2 model.

5.3 E12 — t-SNE/UMAP Feature Space Visualization

Rationale: E8 confirmed via ablation that TransECA-Net learns superior representations compared to CNN-Only. This experiment

uses t-SNE and UMAP to directly visualize whether the deep embeddings produce better-separated clusters than raw features, providing visual confirmation of the representation learning advantage.

Objective: Visualize clustering quality in raw features vs. TransECA-Net learned representations to validate the effectiveness of representation learning.

Method: t-SNE (perplexity=30) + UMAP (n_neighbors=15, min_dist=0.1), executed on both raw features (76-dim) and TransECA-Net embeddings (128-dim), 8,000 stratified subsamples, 13 classes.

Table 22: Clustering Quality (Silhouette Score)

Method	Silhouette Score
t-SNE (Raw)	-0.1621
UMAP (Raw)	-0.2349
t-SNE (Embedding)	-0.0959
UMAP (Embedding)	-0.0668

Table 23: Representation Learning Improvement


Metric	Raw Feature Space	TransECA Embedding
Mean Silhouette	-0.1985	-0.0814
Improvement	—	+59%

Theoretical predictions versus results are compared in Table 21. Silhouette scores are reported in Table 22. The improvement is quantified in Table 23.

Analysis:


- TransECA embeddings improve cluster separation by approximately 59%
- UMAP performs best in embedding space (-0.0668)
- Silhouette scores remaining negative indicate **inherent overlap** among 15 attack classes (especially DoS subclasses) — together with E6’s wide M-F1 CI and E15’s low minority class F1, this points to the same root cause: **inherent similarity between attack subclasses is the systemic bottleneck**
- However, the +59% improvement demonstrates that TransECA-Net learns better discriminative representations than raw features, validating [5]’s core claim about DL’s representation learning advantages

Figure 20: Raw Feature Space (t-SNE)



Note (Figure 20 – Raw Feature Space). The t-SNE projection of 76-dimensional traffic data shows severe inter-class overlap (Silhouette = -0.1985), illustrating why shallow classifiers struggle with fine-grained attack discrimination in the original feature space.

Figure 21: TransECA-Net Embedding Space (UMAP)



Note (Figure 21 — TransECA-Net Embedding Space). The UMAP projection of 128-dimensional embeddings shows improved intra-class cohesion and inter-class separation compared to the raw feature baseline, with the Silhouette Score improving from -0.1985 to -0.0668 (+59%).


Conclusion: The contrast between raw and embedded feature spaces validates [5]’s representation learning hypothesis. TransECA-Net embeddings produce a more separable feature space than the raw input, confirming Stage 2’s value for classifying hard examples.

6 Cross-Experiment Comprehensive Analysis

6.1 Evidence Chain Deduction: System-Level Architecture Performance

The core objective of this project is to resolve the efficiency vs. accuracy trade-off that single models cannot overcome. Having validated Stage 1’s efficiency (E11, E17) and Stage 2’s accuracy (E8) independently, we derive the system-level performance metrics.

Figure 22: The Defense Structures



Note (Figure 22 — The Defense Structures). The hierarchical architecture leverages the orthogonal weaknesses of RF (gradient-immune but noise-fragile) and DL (noise-robust but gradient-susceptible) to achieve resilient defense across multiple attack vectors. System-level metrics are summarized in Table 24.

In-Depth Logic Deduction:

1. Expected Inference Cost ($E[C]$): By the Law of Total Probability, expected time per sample is: $E[C] = C_1 + \alpha(\tau) \cdot C_2$. Substituting measured Stage 1 cost $C_1 = 6.12 \mu s$ (E11), estimated DL cost $C_2 \approx 500 \mu s$ (literature reference [6]; Stage 2 latency benchmarking is identified as future work), and the pass-through rate at optimal tuning $\alpha(\tau = 0.06) = 15.25\%$ (E17):

$$E[C] = 6.12 + 0.1525 \times 500 = \mathbf{82.37 \mu s}$$

Conclusion: Compared to pure deep learning inference on all traffic, we achieved a **6.07× system-level speedup**.

2. System Recall: The probability of detecting an attack is the joint probability of Stage 1 passing it AND Stage 2 correctly classifying it. Assuming independence between the classification errors of the two stages:

$$\begin{aligned} P(\text{Detect}|\text{Attack}) &= P(S_1 = 1|\text{Attack}) \times P(S_2 \text{ correct}|S_1 = 1) \\ &= 0.999 \times 0.93 = \mathbf{92.9\%} \end{aligned}$$

Table 24: System Performance Matrix

Dimension	Metric	Value	Supporting Experiment
Accuracy	System Recall	92.9%	E17 × S2 Training
Accuracy	System FPR	0.007%	E17 × S2 Training
Efficiency	Expected Inference Cost	82.37 μ s	E11 + E17
Efficiency	Speedup (vs DL-Only)	6.07×	E11 + E17
Robustness	System Evasion Rate (PGD $\epsilon=0.01$)	8.04%	E14 + E17
Generalization	UNSW-NB15 W-F1	0.703	E15
Representation	Silhouette Improvement	+59%	E12
Statistics	S2 W-F1 95% CI Width	0.003	E6

Conclusion: Under this independence assumption, we maintain nearly identical detection capabilities as the pure DL-only model (93.0%) at a microscopic fraction of the computational cost. *Limitation:* this independence assumption holds when S1 and S2 errors are uncorrelated; in adversarial scenarios where an attacker jointly optimizes against both stages, actual recall may be lower — the system-level evasion rate analysis (Section 6.1, Point 4) provides the adversarial upper bound.

3. System-Level False Positive Rate (FPR): The upper bound probability of benign traffic raising false alarms is:

$$P(\text{FA}|\text{Normal}) \leq P(S_1 = 1|\text{Normal}) \times P(S_2 = \text{attack}|S_1 = 1) \\ \leq 0.001 \times 0.07 = \mathbf{0.00007 \text{ (or } 0.007\%)}$$

Conclusion: Cascading two filtering mechanisms causes false positive rates to drop multiplicatively, significantly reducing alert fatigue for Security Operation Centers.

4. System Evasion Rate: The final probability of a successful evasion equals the joint probability of bypassing Stage 1 and deceiving Stage 2. Assuming structural independence of vulnerabilities between the two distinct model families (security diversity principle, cf. [15]):

$$P(\text{Successful Evasion}) = \alpha(\tau) \times P(h_2(x + \delta) \text{ is deceived} \mid x \text{ reaches S2})$$

Substituting $\alpha = 15.25\%$, even under the strongest PGD $\epsilon = 0.01$ adversarial attack where h_2 's correct defense rate is 47.28% (deception rate of $1 - 0.4728$):

$$P(\text{Evasion}) = 0.1525 \times (1 - 0.4728) = \mathbf{8.04\%}$$

Conclusion: Even if an advanced attacker compromises the inner Deep Learning layer, the non-differentiable RF front-end limits the system evasion rate to 8.04%.

6.2 Theoretical Models & Statistical Underlying Guarantees

Beyond system-level performance, the evaluation results are supported by statistical theory:

1. Bagging Variance Limits (Breiman's Theorem): According to [11], the generalization error bound of a Random Forest is governed by its base variance and correlation:

$$\text{Var}_{\text{ensemble}} = \rho\sigma^2 + \frac{1-\rho}{B}\sigma^2$$

Our E4 evaluation demonstrated that the OOB Error converges at 0.00171 after $B = 50$ trees. **Conclusion:** As B increases, $\frac{1-\rho}{B}$

approaches zero and residual error is locked to inter-tree correlation ρ . This confirms that $B = 50$ is sufficient for Stage 1.

2. Confidence Interval Behavior Under Extreme Imbalance (Binomial Asymptotics): In E6, extreme minority classes (e.g., Heartbleed with only 3 test samples) exhibited a very wide Bootstrap confidence interval (CI = 0.074). This is not an architectural liability; it stems from the asymptotic variance of multinomial distributions:

$$\text{Var}(\text{M-F1}) \approx \frac{1}{K^2} \sum_{k=1}^K \frac{p_k(1-p_k)}{n_k}$$

Conclusion: CI width is inversely proportional to $\sqrt{n_k}$. When $n_k \rightarrow 3$, variance is inherently large. This confirms that minority class metric instability is an inherent statistical limitation, not a model deficiency.

6.3 Cross-Validation Matrix: Theoretical Predictions vs. Experiments

The full cross-validation matrix is presented in Table 25.

10 verified, 2 overturned. The two overturned predictions do not weaken the hierarchical architecture argument; rather, they reveal more precise mechanisms and the advantage of "orthogonal weaknesses":

- (1) **DL High Variance ([5]) → Overturned:** TransECA-Net showed a Generalization Gap of only +0.006 in E4, indicating low variance. Root Cause: [3]'s observation assumes insufficient data or weak regularization. With 235K training samples and effective regularization (AdamW weight decay + Cosine Annealing + Dropout), variance is suppressed while retaining representational capacity.
- (2) **RF is Naturally Robust to Micro-noise (Bagging Principle) → Overturned:** Under L_∞ uniform noise ($\epsilon = 0.001$) in E14, RF accuracy dropped from 99.59% to 0.65%. This prior belief was based on ensemble stability (Breiman [11] proves variance reduction via bagging but does not claim noise robustness). Root Cause: RF decision boundaries in high dimensions are axis-aligned hyper-rectangles; small uniform noise easily pushes samples across these boundaries. This reveals a fragility of tree-based models to additive noise, despite their immunity to gradient-based attacks.

Strategic System Significance (Orthogonal Weaknesses): If we solely relied on Stage 1 (RF), the system would collapse under

Table 25: Cross-Validation Matrix

Theoretical Prediction	Source	Verification Experiment	Result
Bagging reduces RF Variance	[11]	E4: OOB 50→200 trees nearly unchanged	Verified
DL High Variance	[5]	E4: Gen Gap = 0.006 (low)	Overtured
RF robust to perturbation	Prior belief (ensemble stability)	E14: RF $\epsilon=0.001 \rightarrow 0.65\%$	Overtured
PGD stronger than FGSM	[17]	E14: PGD vs FGSM @ $\epsilon=0.01$: 47% vs 76%	Verified
SHAP axiomatic uniqueness	[16]	E2: SHAP vs Gini $\rho=0.94$	Verified
Learned representations outperform raw features	[5]	E12: Silhouette +59%	Verified
Hierarchical reduces expected cost	Mathematical derivation	E11+E17: 6.07× speedup	Verified
Cross-domain generalization limited by domain distance	[10]	E15: UNSW 64% < CIC 93%	Verified
CI Width $\propto n^{-1/2}$	[13]	E6: S1 Width 0.0002, S2 Width 0.003	Verified
M-F1 CI dominated by minority class samples	$\text{Var} \propto 1/n_k$	E6: M-F1 CI = 0.074 (Heartbleed $n=11$)	Verified

Table 26: Inter-Experiment Cross-Corroboration

Conclusion	Supporting Experiments
Transformer is the architecture core	E8 (Acc +28pp) + E4 (CNN-Only underfitting) + E10 (Attention captures IAT)
ECA benefits minority classes	E8 (M-F1 +0.085) + E10 (ECA CV=0.02 explains W-F1 drop)
Class imbalance is a systemic bottleneck	E6 (M-F1 CI wide) + E12 (Silhouette negative) + E15 (minority class F1 low)
Model decisions align with domain knowledge	E2 (SHAP) + E10 (IG) both focus on Init Win Bytes + IAT
Hierarchical efficiency advantage	E11 (6.12 μ s) + E17 ($\alpha=15.25\%$) $\rightarrow$ 6× speedup
Hierarchical security advantage	E14 (orthogonal weaknesses) + E17 (pass-through rate limitation) $\rightarrow$ evasion rate 8.04%
Performance estimates are reliable	E6 (narrow CI) + E4 (low Variance) + E1 (Nested CV unbiased)

simple noise. If we solely relied on Stage 2 (DL), the system would be vulnerable to PGD gradient attacks and computationally expensive. In the **two-stage architecture**, RF is immune to gradients, and DL resists random noise. Because an attacker cannot craft a payload that is simultaneously unstructured noise and a precise adversarial gradient, the system evasion rate is limited to 8.04%. This is the core justification for the hierarchical design.

6.4 Inter-Experiment Cross-Corroboration

Table 26 maps each conclusion to its supporting experiments.

6.5 Evidence Chain Logic Flow & Closure Graph

Figure 23 presents the complete evidence chain logic flow and closure graph, illustrating how this project progresses from isolated baseline experiments to system-level strategic conclusions across three integrated layers: Stage 1 & 2 verification pools, system-level integration, and global statistical guarantees.


Figure 23: Evidence Chain Logic Flow & Closure Graph



Note (Figure 23 — Evidence Chain Logic Flow & Closure Graph). The figure illustrates the three-layer logical structure of the project:

- (1) **Stage 1 & 2 Verification:** Independent validation of each model's capabilities (E1–E17).
- (2) **System-Level Integration:** Mathematical derivation of system performance ($6.07\times$ speedup, 8.04% evasion rate) from independently measured component metrics.
- (3) **Statistical Guarantee:** Underlying statistical framework (Nested CV, Bootstrap CI) ensuring conclusion reliability.

Figure 24: Hierarchical Architecture Performance (Radar Chart)



Note (Figure 24 — Hierarchical Architecture Performance). The radar chart evaluates the framework across six dimensions: Accuracy, Efficiency, Interpretability, Generalization, Robustness, and Statistical Reliability.

6.5.1 *Detailed Metric Dimension Definitions.* Detailed dimension definitions are given in Table 27.

7 Systemic Bottlenecks and Improvement Directions

7.1 Identified Bottlenecks

The identified bottlenecks are listed in Table 28.

7.2 Improvement Directions

- (1) **Minority class augmentation:** Consider few-shot learning or class-conditional data augmentation for extreme minority classes like Heartbleed/Infiltration
- (2) **Adversarial training:** Introduce adversarial training for TransECA-Net to improve robustness within small perturbation ranges

- (3) **Cross-domain adaptation:** Introduce Domain Adaptation techniques to reduce the feature distribution distance between CIC-IDS2017 and UNSW-NB15
- (4) **Online learning:** Explore incremental learning mechanisms to adapt to temporal drift in network traffic distributions

8 Conclusion

This experiment report comprehensively validates the hierarchical IDS framework through 13 systematic experiments across **6 dimensions**.

8.1 Summary of Core Quantitative Metrics

Table 29 summarizes all core quantitative metrics.

8.2 Comprehensive Dimensional Conclusions

Dimensional conclusions are presented in Table 30.

The hierarchical architecture's core advantage lies not in "each layer being the strongest," but in "**orthogonal weaknesses across two layers + attack surface isolation**": RF filters 84.75% of traffic at $6.12\ \mu\text{s}$ low cost, while TransECA-Net refines the remaining hard examples with DL capability. System performance is statistically guaranteed through Bootstrap CI (narrow) + Nested CV (unbiased) + cross-dataset (UNSW-NB15) triple statistical assurance, ensuring reliable conclusions.

A Experiments Not Executed

Table 31 lists the experiments not executed and their reasons.

B Complete Output File Index

Table 32 provides a complete index of all output files.

References

- [1] I. Sharafaldin, A. H. Lashkari, and A. A. Ghorbani. 2018. Toward generating a new intrusion detection dataset and intrusion traffic characterization. In *Proceedings of the 4th International Conference on Information Systems Security and Privacy (ICISSP)*. 108–116. <https://www.scitepress.org/papers/2018/66398/66398.pdf>
- [2] Z. Liu, Y. Xie, Y. Luo, Y. Wang, and X. Ji. 2025. TransECA-Net: A transformer-based model for encrypted traffic classification. *Applied Sciences* 15, 6 (2025), 2977. <https://doi.org/10.3390/app15062977>
- [3] D. Kwon, H. Kim, J. Kim, S. C. Suh, I. Kim, and K. J. Kim. 2019. A survey of deep learning-based network anomaly detection. *Cluster Computing* 22, Suppl 1 (2019), 949–961. <https://doi.org/10.1007/s10586-017-1117-8>
- [4] M. Ring, S. Wunderlich, D. Scheuring, D. Landes, and A. Hotho. 2019. A survey of network-based intrusion detection data sets. *Computers & Security* 86 (2019), 147–167. <https://doi.org/10.1016/j.cose.2019.06.005>
- [5] J. Wang. 2025. Analysis of machine learning-based methods for network traffic anomaly detection and prediction. In *Proceedings of the 2nd International Conference on Data Science and Engineering (ICDSE)*. 550–554. <https://doi.org/10.5220/0013701800004670>
- [6] Q. Abu Al-Haija, M. Krichen, and W. Abu Elhaija. 2022. Machine-learning-based darknet traffic detection system for IoT applications. *Electronics* 11, 4 (2022), 556. <https://doi.org/10.3390/electronics11040556>
- [7] H. Kaur and G. Singh. 2021. Machine learning techniques for anomaly detection in network traffic. In *Proceedings of the 2021 Sixth International Conference on Image Information Processing (ICIIP)*. 444–449. <https://ieeexplore.ieee.org/document/9702647>
- [8] N. Moustafa and J. Slay. 2015. UNSW-NB15: a comprehensive data set for network intrusion detection systems (UNSW-NB15 network data set). In *Proceedings of the 2015 Military Communications and Information Systems Conference (MilCIS)*. 1–6. <https://ieeexplore.ieee.org/document/7348942>
- [9] S. Abnar and W. Zuidema. 2020. Quantifying attention flow in transformers. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)*. 4190–4197.

Table 27: Radar Chart Evaluation Dimension Definitions

Evaluation Dimension	Core Metrics	Technical Explanation & Physical Meaning
Accuracy	System Recall (92.9%), FPR (0.007%)	Measures the discriminant effectiveness across all 15 attack categories. The hierarchical design ensures high-precision capture of difficult samples after filtering 84.75% of total traffic.
Efficiency	Inference Latency (82 μ s), Speedup (6.07 $\times$)	Measures real-time throughput. Based on the $L_{sys} = L_1 + \alpha L_2$ theoretical derivation, Stage 1's high-speed filtering is proven critical for industrial-scale deployment.
Interpretability (XAI)	SHAP Importance, IG Attribution, Attention Heatmaps	Measures decision transparency. Triangulation via SHAP $\leftrightarrow$ IG proves that the model's decision logic aligns precisely with security domain knowledge like TCP Windows and IAT.
Generalization (Gen)	UNSW-NB15 W-F1 (0.70), Silhouette Coeff (+59%)	Measures adaptability to unseen domains. E15 cross-dataset experiments show that TransECA learns transferable attack feature representations.
Robustness	PGD Evasion Rate (8.04%), Orthogonal Dis-sent	Measures adversarial resilience. The complementary vulnerabilities of non-differentiable trees and deep self-attention reduce the joint system evasion rate.
Statistical Reliability (SR)	Bootstrap CI (± 0.003), Nested CV Gap (0.006)	Measures the certainty of evaluation conclusions. Narrow confidence intervals prove that results are systematic and robust, not isolated random coincidences.

Table 28: Identified Systemic Bottlenecks

Bottleneck	Evidence	Root Cause
Unstable minority class attack classification	E6 M-F1 CI = 0.074; E12 Silhouette < 0	Inherent similarity between attack subclasses + extremely small samples (Heartbleed $n=11$)
Low UNSW-NB15 performance	E15 W-F1 = 0.703 (vs CIC 0.950)	Cross-domain distribution differences (feature space/labeling protocol mismatch)
RF fragile to feature perturbation	E14 $\epsilon=0.001$ causes collapse	Inherent weakness of axis-aligned decision boundaries

Table 29: Summary of Core Quantitative Metrics

Evaluation Domain	Specific Metric	Performance Value	Supporting Experiment
Detection Accuracy	System-Level Global Recall	92.9% (15-class)	E17, S2
Detection Accuracy	System-Level False Positive Rate	0.007%	E17, S2
Processing Performance	Avg. Inference Latency per Sample	82.37 μ s	E11, E17
Processing Performance	Relative Speedup (vs Monolithic DL)	6.07 $\times$	E11, E17
Adversarial Resilience	System Evasion Rate under Strongest PGD	8.04%	E14, E17
Base Capability (S1)	Stage 1 Filter Recall (Binary)	99.9%	E16, E1
Base Capability (S2)	Stage 2 Multiclass Weighted F1	0.950	S2 Training
Generalization Ability	UNSW-NB15 Cross-Dataset W-F1	0.703	E15
Rep. Learning	Silhouette Coeff. Gain in UMAP Space	+59%	E12
Statistical Confidence	W-F1 95% Confidence Interval (CI) Width	0.003	E6

Table 30: Comprehensive Dimensional Conclusions

Dimension	Core Conclusion	Key Data
(1) Accuracy	System Recall 92.9%, FPR < 0.01%	E17 × S2
(2) Efficiency	6.07× speedup, 82 μs/sample	E11 + E17
(3) Interpretability	SHAP ↔ IG triangulation, decisions align with domain knowledge	E2 + E10
(4) Generalization	UNSW W-F1=0.70, Silhouette +59%	E15 + E12
(5) Robustness	Orthogonal weaknesses, system evasion rate 8.04%	E14
(6) Statistical Reliability	CI ∝ $n^{-1/2}$, Nested CV unbiased	E6 + E1

Table 31: Appendix A: Experiments Not Executed

Experiment	Reason
E3 (Feature Selection Top-K)	RF has built-in feature importance; E2 SHAP already covers this with marginal additional benefit
E5 (Validation Curve)	E1 Nested CV grid search already includes hyperparameter vs. performance curves; conclusions fully covered
E7 (DL Regularization)	E8 ablation already validates architectural contributions; regularization tuning overlaps with E8
E13 (Imbalance Handling SMOTE)	class_weight already included in E1’s search space; merged into execution

Table 32: Appendix B: Complete Output File Index

Experiment	Core Output
E1+E16	stage1_rf_best.pkl
E17	results/E17_threshold_tuning*.json
E11	results/E11_latency_benchmark*.json
S1	models_chk/stage1_rf_stratified.joblib, data/stage2/*.parquet
S2	models_chk/stage2_transeca.pth
E8	results/E8_ablation_comparison.png, results/E8_ablation_bar.png
E15	results/E15_unsw_training_curves.png, results/E15_unsw_confusion_matrix.png, results/E15_cross_dataset_comparison.png
E2	results/E2_shap_summary.png, results/E2_shap_bar.png, results/E2_shap_vs_rf.png
E6	results/E6_bootstrap_distributions.png
E10	results/E10_ig_global_importance.png, results/E10_attention_heatmap.png, results/E10_eca_channel_weights.png
E14	results/E14_robustness_curve.png, results/E14_per_class_robustness.png
E4	results/E4_oob_vs_trees.png, results/E4_dl_learning_curves.png, results/E4_complexity_vs_perf.png, results/E4_rf_learning_curve.png
E12	results/E12_tsne_raw.png, results/E12_umap_raw.png, results/E12_tsne_embedding.png, results/E12_umap_embedding.png

- [10] S. Ben-David, J. Blitzer, K. Crammer, A. Kulesza, F. Pereira, and J. W. Vaughan. 2010. A theory of learning from different domains. *Machine Learning* 79, 1-2 (2010), 151–175. <https://doi.org/10.1007/s10994-009-5152-4>
- [11] L. Breiman. 2001. Random forests. *Machine Learning* 45, 1 (2001), 5–32. <https://doi.org/10.1023/A:1010933404324>
- [12] G. C. Cawley and N. L. C. Talbot. 2010. On over-fitting in model selection and subsequent selection bias in performance evaluation. *Journal of Machine Learning Research* 11 (2010), 2079–2107. <https://jmlr.org/papers/v11/cawley10a.html>
- [13] B. Efron. 1979. Bootstrap methods: another look at the jackknife. *The Annals of Statistics* 7, 1 (1979), 1–26. <https://doi.org/10.1214/aos/1176344552>
- [14] I. J. Goodfellow, J. Shlens, and C. Szegedy. 2015. Explaining and harnessing adversarial examples. In *International Conference on Learning Representations (ICLR)*.
- [15] B. Littlewood and L. Strigini. 2004. Redundancy and diversity in security. *IEEE Security & Privacy* 2, 3 (2004), 56–61.
- [16] S. M. Lundberg and S. I. Lee. 2017. A unified approach to interpreting model predictions. In *Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 30. <https://proceedings.neurips.cc/paper/2017/hash/8a20a8621978632d76c43dfd28b67767-Abstract.html>
- [17] A. Madry, A. Makelov, L. Schmidt, D. Tsipras, and A. Vladu. 2018. Towards deep learning models resistant to adversarial attacks. In *International Conference on Learning Representations (ICLR)*. <https://arxiv.org/abs/1706.06083>
- [18] Q. Wang, B. Wu, P. Zhu, P. Li, W. Zuo, and Q. Hu. 2020. ECA-Net: Efficient channel attention for deep convolutional neural networks. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 11534–11542.